

README

Project Information

Authors:

- Victor Allen, Vince Gonzales

Username:

- Victor Allen: cssc1404@edoras.sdsu.edu (Use this account for testing)
- Vince Gonzales: cssc1422@edoras.sdsu.edu

Class: CS 480 – Summer 2025

Assignment: Assignment #3 – Memory Simulation

Files Included (File Manifest):

- |— main.cpp
- |— generate_graphs.py
- |— Makefile – builds the executable "sim"
- |— Simulator.h
- |— Simulator.cpp
- |— MemoryManager.h
- |— MemoryManager.cpp
- |— MemoryBlock.h
- |— MemoryBlock.cpp
- |— README – this file

Compile Instructions:

1. Open a terminal and navigate to your assignment directory:
`cd ~/a3`
2. Run make:
`make`

This produces the executable `sim`.

Operating Instructions:

1. Run the simulation:
`./sim`
2. To generate graphs after running the simulation:
`python3 generate_graphs.py`

Note: Requires matplotlib, pandas, and numpy. Install with:
`pip3 install matplotlib pandas numpy`

Design Decisions:

The simulation implements two separate MemoryManager instances running identical request sequences to ensure fair comparison between First Fit and Best Fit algorithms. Key design choices include:

- **Linked List Implementation:** Memory blocks are managed using a singly-linked list where each node represents a contiguous memory region (either allocated or free)
- **Synchronization Strategy:** Both algorithms process identical request sequences to eliminate randomness as a factor in performance differences
- **Fragment Definition:** External fragments are defined as free blocks of size 1 or 2 units, representing unusable memory holes
- **Request Generation:** 50/50 split between allocation and deallocation requests with random sizing between 3-10 units
- **Performance Tracking:** Statistics are updated after each request and periodically saved for time-series analysis


Lessons Learned:

- **Algorithm Trade-offs:** The simulation clearly demonstrates the classic trade-off between allocation speed (First Fit advantage) and memory utilization efficiency (Best Fit advantage)
- **Fragmentation Patterns:** Both algorithms eventually reach steady-state fragmentation levels, but Best Fit consistently maintains lower fragmentation
- **Implementation Complexity:** Maintaining synchronization between two parallel memory managers required careful error handling and rollback mechanisms
- **Statistical Significance:** Large-scale simulation (10,000 requests) was necessary to observe meaningful performance differences and achieve statistical stability

Graphs:


The simulation generates three comprehensive visualizations:

1. Performance Comparison Bar Chart




2. Detailed Performance Analysis

Detailed Performance Analysis: First Fit vs Best Fit



3. Performance Evolution Over Time



Background:

This simulation compares two fundamental memory allocation algorithms in operating systems:

First Fit Algorithm:

- Searches linked list sequentially from the beginning
- Allocates the first block large enough to satisfy the request
- Time complexity: $O(n)$ average case, $O(1)$ best case
- Strategy: Speed over optimization

Best Fit Algorithm:

- Searches the entire linked list to find the smallest suitable block
- Always minimizes wasted space in each allocation
- Time complexity: $O(n)$ worst case (complete traversal required)
- Strategy: Memory efficiency over speed

Simulation Parameters:

- Memory size: 256 KB (128 units \times 2 KB each)
- Total requests: 10,000
- Request sizes: 3-10 units (randomly distributed)
- Request pattern: 50% allocation, 50% deallocation
- Fragmentation metric: Count of holes sized 1-2 units

Performance Metrics:

1. **Average External Fragments per Request:** Measures memory fragmentation
2. **Average Nodes Traversed per Allocation:** Measures computational overhead
3. **Percentage Allocation Requests Denied:** Measures algorithm success rate

Findings:

The simulation results reveal significant performance characteristics:

Fragmentation Management:

- Best Fit produces substantially fewer external fragments, typically 20-30% less than First Fit
- Both algorithms reach fragmentation steady-state after \sim 5,000-7,000 requests
- First Fit creates clustering of small fragments at the beginning of memory

Computational Efficiency:

- First Fit demonstrates superior search efficiency, averaging 60-70% fewer node traversals
- Best Fit's exhaustive search requirement creates consistent $O(n)$ overhead
- Performance gap widens as fragmentation increases list length

Memory Utilization:

- Both algorithms show similar denial rates (typically 1-5%), indicating that algorithm choice is less critical than total memory capacity
- Best Fit's improved space utilization provides marginal improvement in allocation success
- High denial rates suggest the 256KB memory size may be insufficient for the request pattern